

SQL – Data Manipulation Language Part 2

Lecturer: Neena Thota

neena.thota@it.uu.se

Intended Learning Outcomes


- Write and understand:
 - Nested queries.
 - Queries based on views.
 - Outer and Inner Joins.

Review

```
select target_list  
from table_list  
[ where tuple_predicates ]  
[ group by attribute_list]  
[ having group_predicates]  
[ order by attribute_list + ASC/DESC ]
```

NESTED QUERIES

Nested Queries



- Require that existing values in database be fetched and then used in a comparison condition.
- **E.g.** Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

General interpretation of Nested Queries

SELECT * ← External query
 FROM Relation R ← reference
 WHERE Attribute in (SELECT ...
) ← Nested query (NQ)
↑ predicate → Execute NQ, then test predicate

SELECT
 FROM
 WHERE

DISTINCT Pnumber

PROJECT

Pnumber IN

(
 SELECT
 FROM
 WHERE

Pnumber

PROJECT, DEPARTMENT, EMPLOYEE

Dnum=Dnumber AND

Mgr_ssn=Ssn AND Lname='Smith')

IN (comparison operator)

OR

Pnumber IN

(
 SELECT
 FROM
 WHERE

Pno

WORKS_ON, EMPLOYEE

Essn=Ssn AND Lname='Smith');

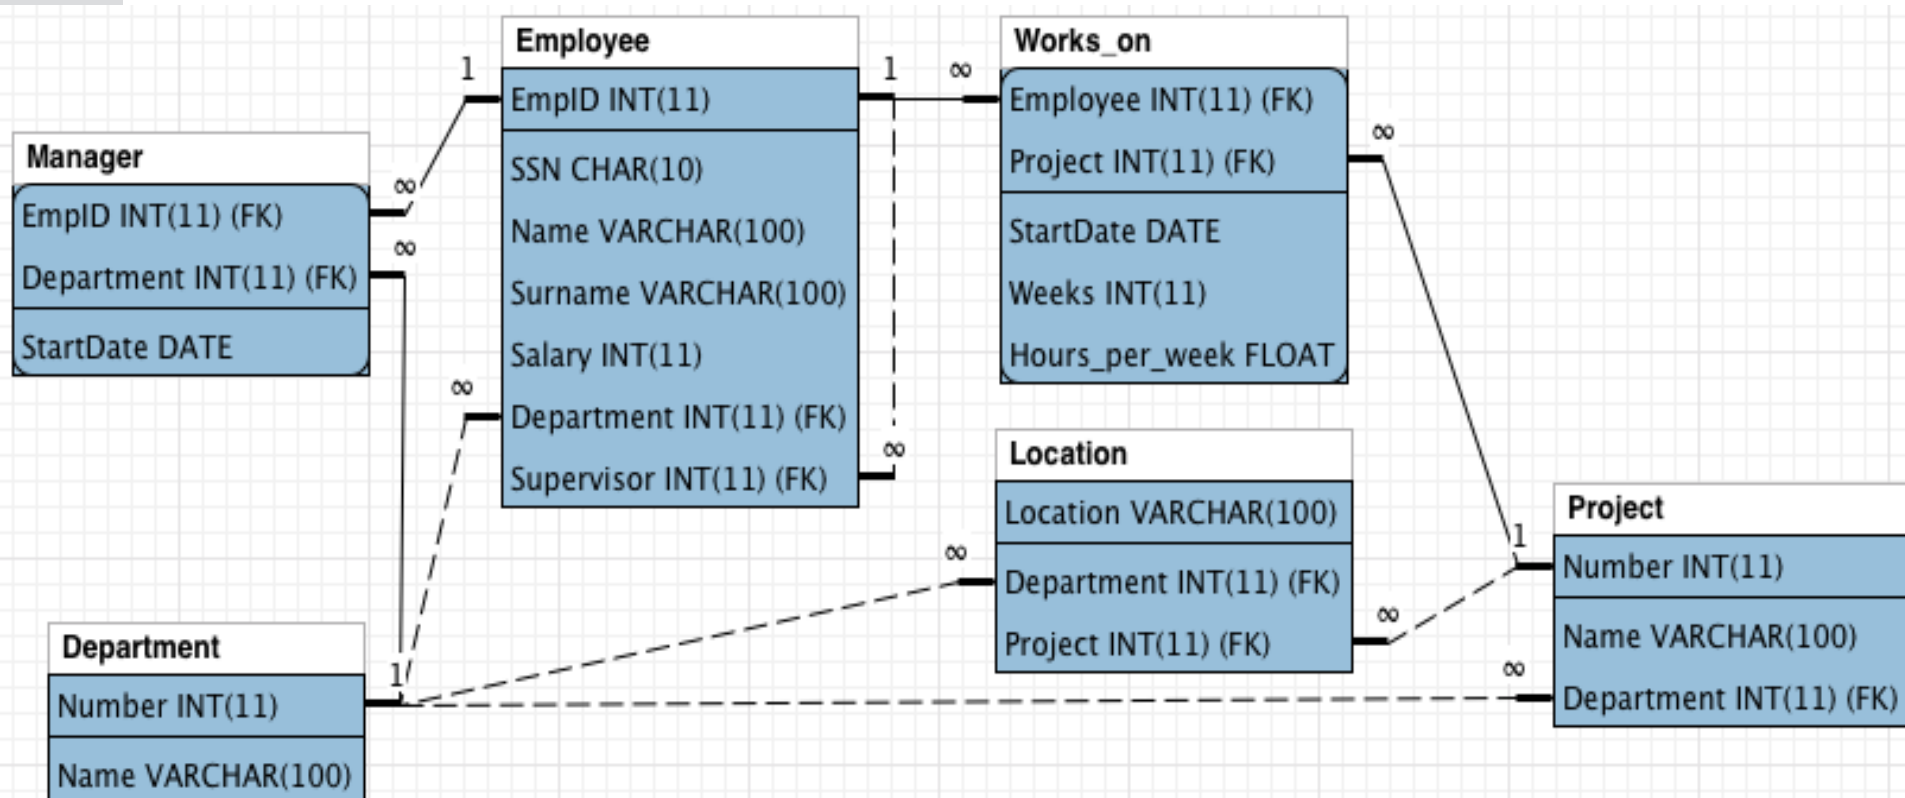
OR (logical operator)

Operators for Nested Queries

- When the internal (nested) query returns more than one tuple, we need to better specify the external predicate.
- We can use normal comparison operators like $<$, $>$, $>=$, $=$, etc. followed by the operators **ANY** and **ALL**. For example:
 - $=$ ANY (can also be written: **IN**)
 - $<>$ ALL (can also be written: **NOT IN**)
- **EXISTS** checks if the result of an internal query is empty.
 - The opposite is: **NOT EXISTS**.

EXAMPLES

1. Select all the employees whose supervisor is a manager (using the **=ANY** or **IN** operators).
2. Select all the departments where there are no employees called “Doo” (using the **<>ALL** or **NOT IN** operators).
3. Select all the employees having the maximum salary with respect to their department (using the **NOT EXISTS** operator).



E.g.1. Select all the employees whose supervisor is a manager (using ANY or IN).

```
select *  
from Employee  
where Supervisor = ANY (Select EmpID  
                        from Manager)
```

-- or

```
select *  
from Employee  
where Supervisor IN (Select EmpID  
                    from Manager)
```

E.g.2. Select all the departments where there are no employees called “Doo” (using the ∇ *ALL* or *NOT IN*).

```
select *  
from Department  
where Number  $\nabla$ ALL (select Department  
                        from Employee  
                        where Surname = ‘Doo’)
```

```
-- or  
select *  
from Department  
where Number NOT IN (select Department  
                      from Employee  
                      where Surname = ‘Doo’)
```

E.g.3. Select all the employees having the maximum salary with respect to their department (using *NOT EXISTS*)

```
select *  
from Employee E  
where NOT EXISTS (select *  
                    from Employee  
                    where Department = E.Department  
                    and Salary > E.Salary)
```

TO COMPARE SINGLE ATTRIBUTE

Summary

- $= \text{ANY}$ (or $= \text{SOME}$) operator returns TRUE if value v is equal to *SOME* value in set V and is hence equivalent to IN.
- Comparison condition $(v > \text{ALL } V)$ returns TRUE if value v is greater than *ALL* values in set (or multiset) V .
- Result of EXISTS is Boolean value TRUE if nested query result contains at least 1 tuple, or FALSE if nested query result contains no tuples.

Visibility

- Each relation/alias is visible only inside internal subqueries.
- Rule: reference to an *unqualified attribute* refers to the relation declared in the **innermost nested query**.
- Advisable to create tuple variables (aliases) for *all the tables referenced in an SQL query*

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------

```

SELECT      E.Fname, E.Lname
FROM        EMPLOYEE AS E
WHERE       E.Ssn IN ( SELECT      Essn
                        FROM        DEPENDENT AS D
                        WHERE       E.Fname=D.Dependent_name
                        AND E.Sex=D.Sex );
  
```



VIEWS

Views (Virtual Tables)

- **View** – a single (virtual) table derived from other tables; can be from **base tables (physical form)** or other **previously defined views**.
- To create a view:

CREATE VIEW *viewname* **AS** <query expression>

CREATE VIEW RichEmployees AS

select *
from Employee
where Salary > 60000

CREATE VIEW ProjectWork(Emp, Project) AS

select E.Surname, P.Name
from Employee E, Works_on W, Project P
where E.EmpID=W.Employee AND P.Number=W.Project

**Specifies new
attribute names**



Update of a view

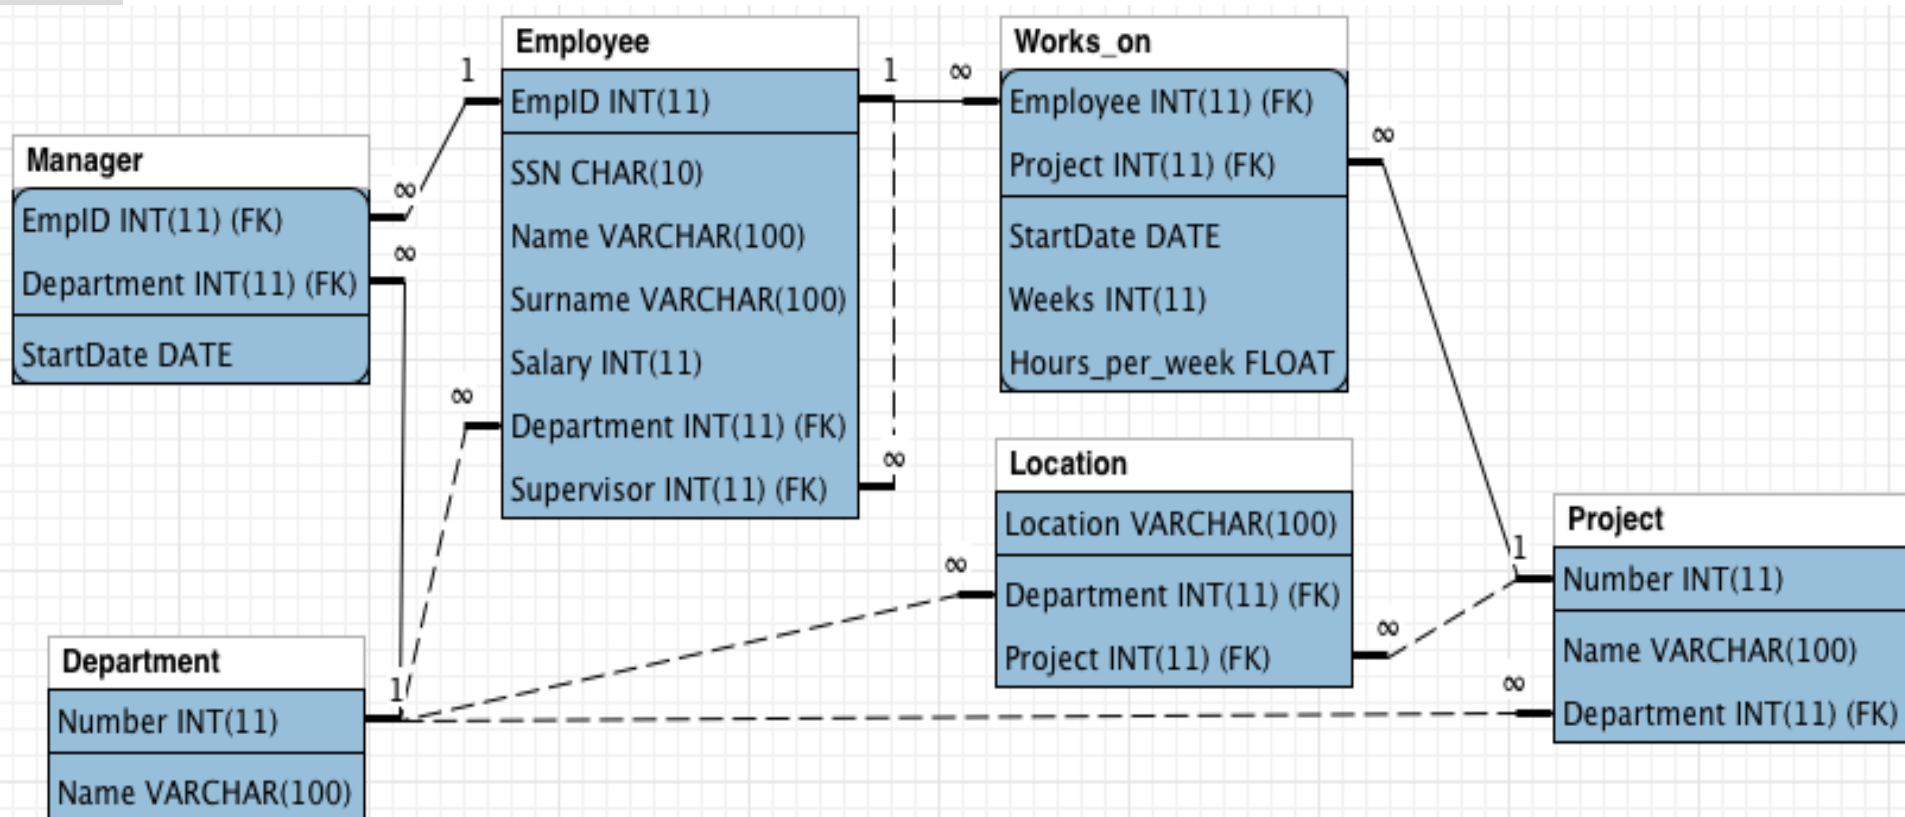
DBMS

- Can modify or transform view query (submitted by user) into query on underlying base tables – **Query Modification**.
- Physically creates temporary view table when view is first queried and keeps table on assumption that other queries on view will follow – **View Materialization**.
- View/s
 - with a single definition table is **updatable** if its attributes **contain the PK**.
 - on multiple tables **using joins** are **not updatable** in general.
 - using **GROUP BY** and **aggregations** are **not updatable**.
- Advanced options can be specified using SQL for security and authorization (not treated here).
- We can use the **DROP VIEW** command to dispose of a view.

Exercise

- Express the following query defining a view (or more than one, if you prefer), and writing a query that uses the view as a base relation.

Select the department(s) with the highest average salary.



Solution, previous slide

Select the department(s) with the highest average salary.

```
CREATE VIEW dept_avg_salary AS  
SELECT Department, AVG(Salary) as avg_salary  
FROM Employee  
GROUP BY Department;
```

```
SELECT Department  
FROM dept_avg_salary  
WHERE avg_salary >= ALL (Select avg_salary  
from dept_avg_salary);
```



JOINS

Some Types of Joins

- **NATURAL JOIN**
 - no join condition is specified; implicit *condition* for *each pair of attributes with same name* from *both tables* is created.
 - Each such pair of attributes included *only once* in resulting relation
- **INNER JOIN**
 - only pairs of tuples that **match join condition** are retrieved,
- **OUTER JOIN**
 - tuples from two tables **combined** by matching corresponding rows without losing any tuples for lack of matching values.
- **LEFT OUTER JOIN**
 - every tuple in the **left** table must appear in the result;
 - If no matching tuple, padded with NULL values for attributes of right table.
- **RIGHT OUTER JOIN**
 - Every tuple in the **right** table must appear in the result;
 - If no matching tuple, padded with NULL values for attributes of left table.
- **FULL OUTER JOIN**
 - Keeps all tuples in both **left and right relations** when no matching tuples are found, padding with NULL values as needed.

EXAMPLES

<i>A</i>	<i>B</i>
1	2
3	4

(a) Relation *R*

<i>B</i>	<i>C</i>	<i>D</i>
2	5	6
4	7	8
9	10	11

(b) Relation *S*

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
1	2	5	6
3	4	7	8

NATURAL JOIN

<i>A</i>	<i>B</i>	<i>C</i>
1	2	3
4	5	6
7	8	9

(a) Relation *U*

<i>B</i>	<i>C</i>	<i>D</i>
2	3	10
2	3	11
6	7	12

(b) Relation *V*

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
1	2	3	10
1	2	3	11
4	5	6	⊥
7	8	9	⊥
⊥	6	7	12

OUTER JOIN

LEFT JOIN

RIGHT JOIN

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
1	2	3	10
1	2	3	11
4	5	6	⊥
7	8	9	⊥

<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>
1	2	3	10
1	2	3	11
⊥	6	7	12

EXAMPLE

```
select EmpID, Name  
from Manager right outer join Department  
on Department = Number
```

- Using a simple join between managers and departments, a department without a manager is not selected.
- We can extend the result by including those departments that do not match any manager.
 - (in general: with those tuples not matching any tuple in the other relation).
- The rest of the tuple in the result is filled with *nulls*.

Example of right outer join

```
select EmpID, Name
from Manager right outer join Department
on Department = Number
```

MANAGER

<i>EmpID</i>	<i>Department</i>
1	45
2	48

DEPARTMENT

<i>Number</i>	<i>Name</i>
45	IT
48	Phyl
51	Math

Result:

<i>EmpID</i>	<i>Name</i>
1	IT
2	Phyl
<i>NULL</i>	Math

Similarly, we can use a **left outer join** or a **full outer join**.

Wrap up: SQL as a DML

- Simple SQL queries.
- Extensions of SELECT-FROM-WHERE.
 - LIKE, BETWEEN, IS (NOT) NULL.
 - DISTINCT.
 - ORDER BY.
- Set operators
 - UNION, (EXCEPT, INTERSECT)
- Aggregation
 - GROUP BY / HAVING.
- Nested queries
 - ANY, ALL, (NOT) IN, (NOT) EXISTS.
- Views.
- Outer and Inner joins.