



UPPSALA
UNIVERSITET

Learning from nature

Olle Gällmo

Uppsala University
Department of Information Technology



Natural computation

- Using computers to model/simulate natural phenomena
 - ✱ to learn more about these phenomena
 - ✱ to learn new ways to solve computational problems
 - ✱ to learn how to build computational devices from biological material

Computers and humans




What is learning?

- The ability to improve over time, based on experience
- Why?
 - ✱ Solutions to problems are not always programmable
- Examples
 - ✱ Handwritten character recognition
 - ✱ Adaptive control of production processes
 - ✱ Game programs that adjust parameters and/or strategies over time
 - ✱ Learning to walk by trial-and-error




Three forms of learning


Supervised



Reinforcement



Unsupervised





Techniques (examples)

- Artificial neural networks (ANNs)
 - ✱ Inspired by biological nervous systems
 - ✱ E.g. Multilayer perceptrons, Self-Organizing Maps
- Reinforcement learning (RL)
 - ✱ Inspired by psychology, ethology and behaviourism
 - ✱ E.g. Menace, Q-Learning, $TD(\lambda)$
- Evolutionary Computing (EC)
 - ✱ Inspired by genetics, natural selection and evolution
 - ✱ E.g. Genetic algorithms, Genetic Programming
- Swarm intelligence
 - ✱ Inspired by social animals (bird flocks, ants, etc.)
 - ✱ E.g. Particle Swarm Optimization, Ant Colony Optimization, Cellular automata



Techniques (examples)

■ Artificial neural networks (ANNs)

- ✱ Inspired by biological nervous systems
- ✱ E.g. Multilayer perceptrons, Self-Organizing Maps

■ Reinforcement learning (RL)

- ✱ Inspired by psychology, ethology and behaviourism
- ✱ E.g. Menace, Q-Learning, TD(λ)


■ Evolutionary Computing (EC)

- ✱ Inspired by genetics, natural selection and evolution
- ✱ E.g. Genetic algorithms, Genetic Programming

■ Swarm intelligence

- ✱ Inspired by social animals (bird flocks, ants, etc.)
- ✱ E.g. Particle Swarm Optimization, Ant Colony Optimization, Cellular automata


Artificial neural networks



- Begun in the 1940's
- Many simple processing elements (neurons), operating in parallel and communicating through weighted connections
- Based on very simple models of biological neurons and synaptic connections
- Used both for industrial applications and as a model to study biological systems




An artificial neuron




$$y = f(S)$$

$$S = \sum_{i=1}^n w_i x_i - \theta = \sum_{i=0}^n w_i x_i$$

$f(S)$ = any non-linear, saturating function, e.g. a step function or a sigmoid:



Multilayer perceptrons

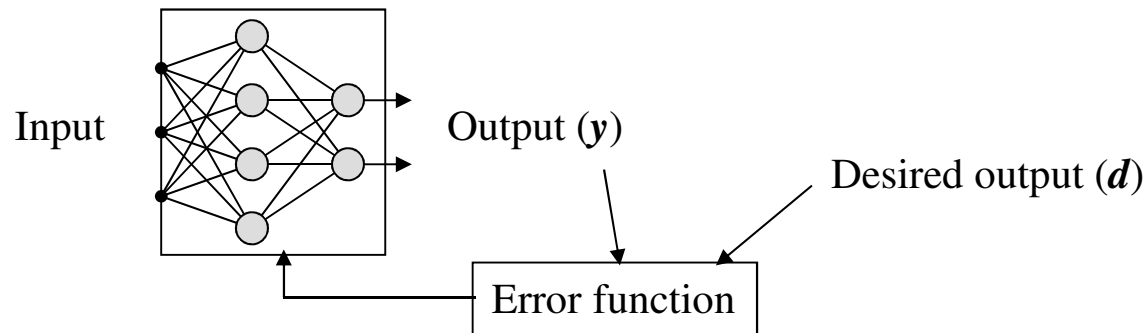


Can approximate any function to any degree of accuracy, given a sufficiently rich internal structure (number of nodes and layers)

Most common training algorithm: *Back propagation*




Back propagation



The contribution to the error E from a particular weight w_{ji} is $\frac{\partial E}{\partial w_{ji}}$

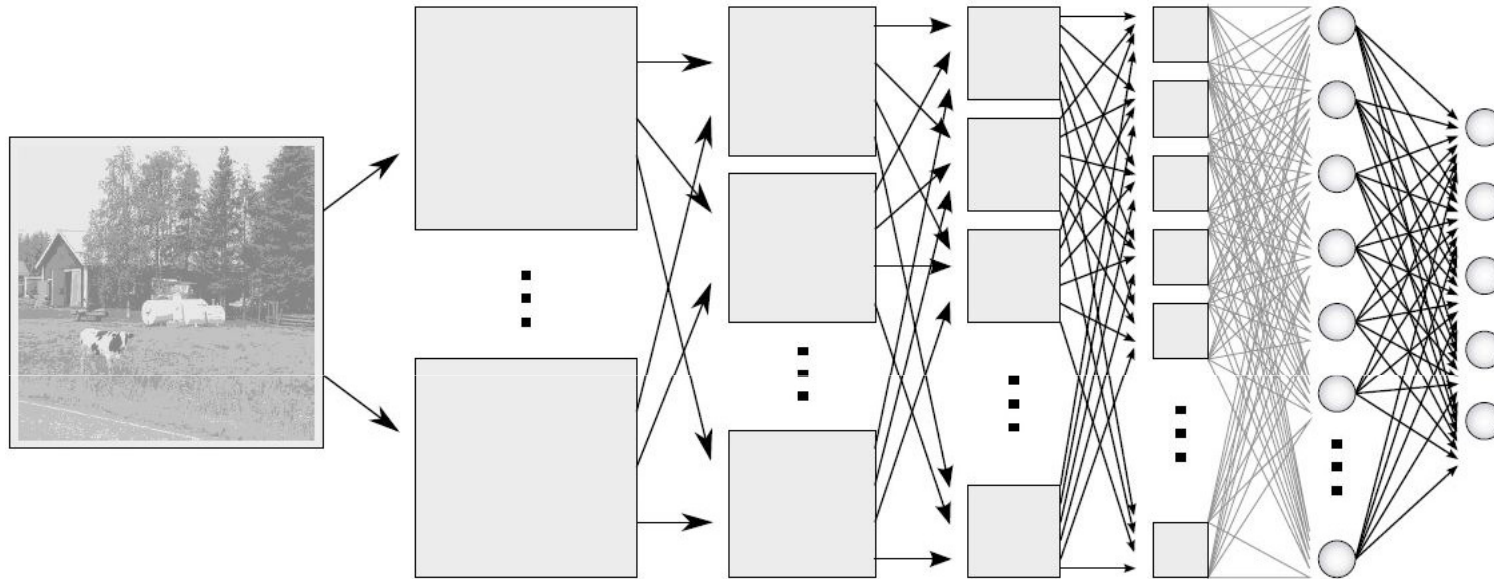
The weight should be moved in proportion to that contribution, but in the other direction:

$$\Delta w_{ji} = -\eta \frac{\partial E}{\partial w_{ji}}$$



Error function and activation function must both be differentiable.

Deep Learning



- A very deep multilayer structure (many layers)
- Back propagation ineffective for that many layers
- Layers instead trained separately

Artificial neural networks ...

- store information in the weights, not in the nodes
- are trained, by adjusting the weights, not programmed
- can generalize to previously unseen data
- are adaptive
- are concurrent
 - ✱ well suited for parallel simulation and/or hardware implementation
- are fault tolerant



UPPSALA
UNIVERSITET

Early neurocomputers






Fig. 1.5 • Charles Wickens, 1960. The first 60 test-tube pairs. Far



Adam




Techniques (examples)

- Artificial neural networks (ANNs)
 - ✱ Inspired by biological nervous systems
 - ✱ E.g. Multilayer perceptrons, Self-Organizing Maps
- **Reinforcement learning (RL)**
 - ✱ Inspired by psychology, ethology and behaviourism
 - ✱ E.g. Menace, Q-Learning, TD(λ)
- Evolutionary Computing (EC)
 - ✱ Inspired by genetics, natural selection and evolution
 - ✱ E.g. Genetic algorithms, Genetic Programming
- Swarm intelligence
 - ✱ Inspired by social animals (bird flocks, ants, etc.)
 - ✱ E.g. Particle Swarm Optimization, Ant Colony Optimization, Cellular automata



Reinforcement learning



- **Reward:** an evaluation of the environmental state (only indirectly an evaluation of the agent's actions)
- **Goal:** To make decisions (find actions) that maximise the long term reward received by the agent.
- The agent must be allowed to *explore*, i.e. sometimes do actions that at the time seem sub-optimal.
- Learning by trial-and-error



UPPSALA
UNIVERSITET

MENACE

(D. Michie 1961)




Plate 1. The original matchbox version of MENACE

Typical RL problems

- Games
- Autonomous robots
- Control of unstable systems
 - ✱ Learning to ride a bicycle
 - ✱ Auto-pilot for helicopters
- Sequential optimization problems, for example:
 - ✱ Controlling the elevators in an office building
 - ✱ Resource allocation in computer networks



Techniques (examples)

- Artificial neural networks (ANNs)
 - ✱ Inspired by biological nervous systems
 - ✱ E.g. Multilayer perceptrons, Self-Organizing Maps
- Reinforcement learning (RL)
 - ✱ Inspired by psychology, ethology and behaviourism
 - ✱ E.g. Menace, Q-Learning, TD(λ)
- **Evolutionary Computing (EC)**
 - ✱ Inspired by genetics, natural selection and evolution
 - ✱ E.g. Genetic algorithms, Genetic Programming
- Swarm intelligence
 - ✱ Inspired by social animals (bird flocks, ants, etc.)
 - ✱ E.g. Particle Swarm Optimization, Ant Colony Optimization, Cellular automata

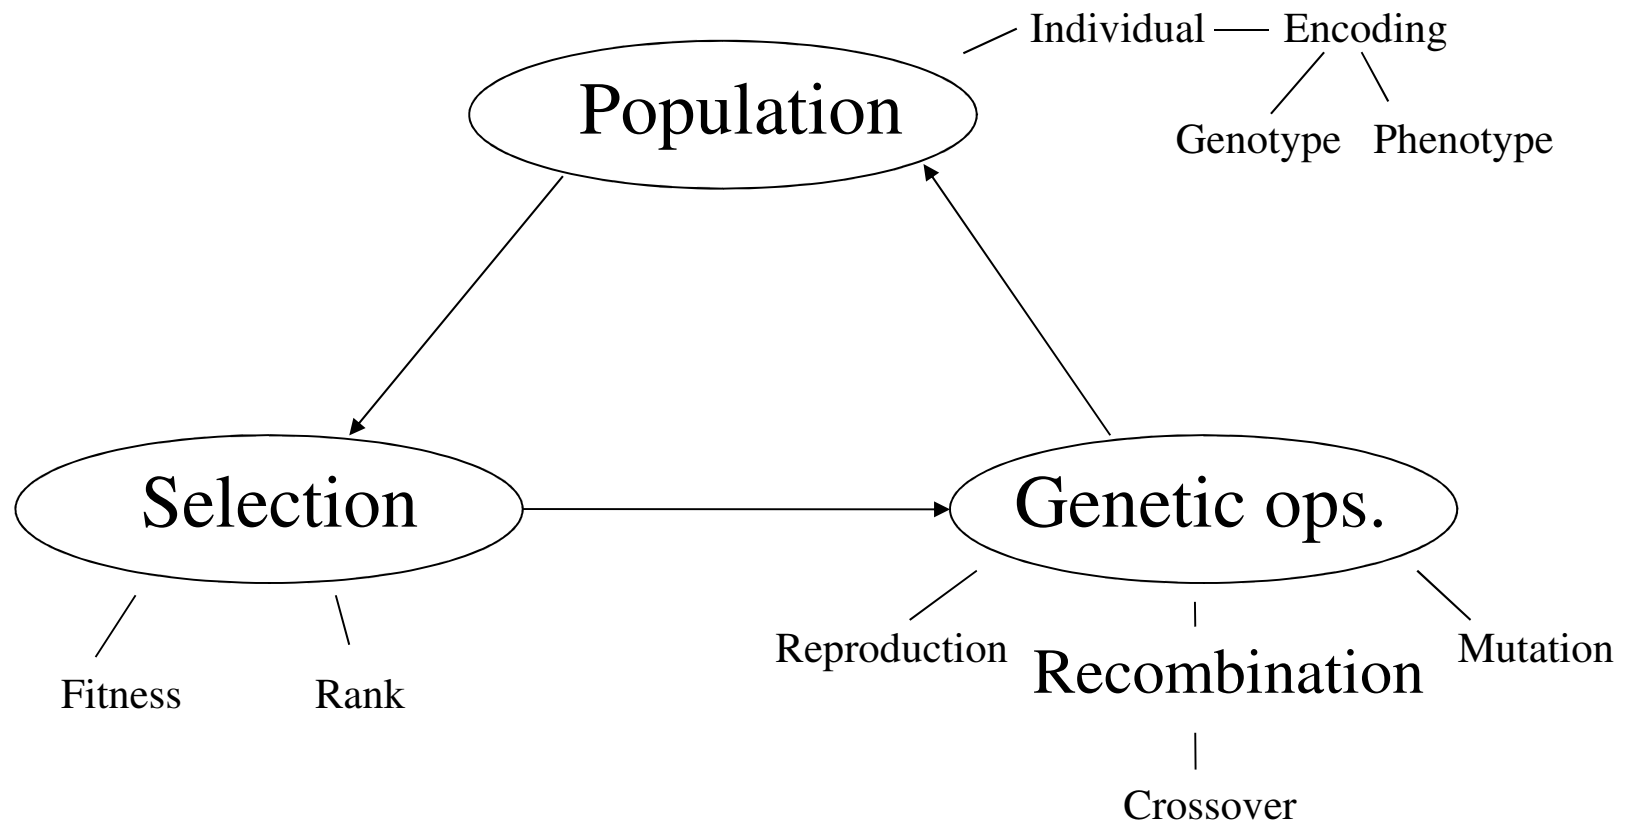


Evolutionary computing

- Used for learning problems where the task is to maximize some measure of success (fitness)
- Essentially the same family of *problems* as in reinforcement learning, but the *methods* are different
- Methods inspired by genetics, natural selection and evolution
- However, the "evolution" is controlled, so it's more like breeding



Evolutionary computing



Genotypes

- A solution to the problem is encoded by the individual's *genotype* (*genome*, *artificial chromosome*)
 - ✱ In genetic algorithms, a string or parameter vector (e.g. bit string)
 - ✱ In genetic programming, a computer program
 - ✱ In evolutionary programming, a representation of a state machine
 - ✱ ...


Genetic Programming

- Usually operates on parse trees of computer programs
- E.g. the expression $5+3*4$

In Lisp

`(+ 5 (* 3 4))`


As a tree



- Crossover: Swap sub trees
- Works particularly well in Lisp!

GP example: function approximation


- Task: Given a training set, discover the function



- A neural network would do a *numerical* approximation
- GP is *combinatorial* – it should be able to find the *exact* function (if given the necessary building blocks)




Implementation



- Create a population of random expressions, $y(x)$, using the functions $+$, $-$, $*$, $/$, \sin , \cos , \exp and \log , and the terminals 1 and x
- Many lures (\exp in particular)
- Fitness: 0 if illegal expression, else $1/(d+1)$, where $d = |f(x) - y(x)|$
- This way, fitness stays in $[0,1]$

Test run (100 individuals)

Found functions





Techniques (examples)


- Artificial neural networks (ANNs)
 - ✱ Inspired by biological nervous systems
 - ✱ E.g. Multilayer perceptrons, Self-Organizing Maps
- Reinforcement learning (RL)
 - ✱ Inspired by psychology, ethology and behaviourism
 - ✱ E.g. Menace, Q-Learning, TD(λ)
- Evolutionary Computing (EC)
 - ✱ Inspired by genetics, natural selection and evolution
 - ✱ E.g. Genetic algorithms, Genetic Programming
- **Swarm intelligence**
 - ✱ Inspired by social animals (bird flocks, ants, etc.)
 - ✱ E.g. Particle Swarm Optimization, Ant Colony Optimization, Cellular automata

Swarm Intelligence

- Bird flocks and fish schools move in a coordinated way, but there is no coordinator (leader)
 - ✱ So, what decides the behaviour of a leader-less flock?
- Ants and termites quickly find the shortest path between the nest and a food source
 - ✱ ... and solve many other advanced problems as well
 - keeping cattle, building (ventilated) housing, coordinated heavy transports, tactical warfare, cleaning house, etc.
 - ✱ A single ant is essentially a blind, memory-less, random walker!
- Distributed systems without central control
- Useful not only to simulate but also to solve optimization problems

Bird flocks and fish schools

- Local interaction
- No leader
- Simple local rules – a weighted combination of several goals
 - ✱ match velocity of your neighbours
 - ✱ avoid collisions with your neighbours
 - ✱ avoid getting too far from your neighbours
 - or strive for centre of the flock (fish)
- To simulate an insect swarm, remove the match-velocity rule
- Sufficient to make very realistic simulations of fish schools and bird flocks
 - ✱ used in movies and computer graphics





UPPSALA
UNIVERSITET


Stampede in "Lion King"




Particle Swarm Optimization

Lbest with $\varphi_1=1.8$, $\varphi_2=2.3$

- Nhood: the 2 immediate neighbours
- $V_{max} = \text{range}/25$



Bowl



Alpine 2D




What about the ants?


- How do they find the shortest route?
 - ✱ They don't (not the individual ants, that is)
 - ✱ The colony does!
- Ant colonies are much more intelligent than ants
 - ✱ Ant colonies adapt, ants don't (much)
 - ✱ Ants have almost no memory and can not build cognitive maps. Ant colonies can (and do)
 - Mammals build cognitive maps in their brains
 - Ant colonies build them in their environment, through pheromone trails
- Ants are better thought of as cells in a greater organism – the colony
 - ✱ Also without leader – the queen is not a controller




Ants find shortest paths



At first, the ants
select at random



After a while, pheromones
become more concentrated
on the shortest route



Drawings by Marco Dorigo

Stigmergy

Indirect communication and coordination, by local modification and sensing of the environment





Ant Colony Optimization


- Family of combinatorial optimization algorithms, based on ant behaviour
- Common benchmark: the Travelling Salesman Problem (TSP)
- Common 'real' applications
 - ✱ Scheduling and
 - ✱ Network routing (AntNet)
- Members: ACS, Ant-Q, MMAS, AS_{rank} , ...
 - ✱ most of which are extensions to Dorigo's Ant System (AS)

Cellular automata

- Massively parallel system of identical communicating state machines (cells)
- A cell's *state* (e.g. on/off) is a function of the states of it communicates with (its neighbours)
 - ✱ The neighbourhood is usually topological
- Used to model/animate fluids (*Find Nemo*), gases, bacterial growth, swaying grass (*Shreck?*), social interaction, epidemics, in ecological simulations etc.

Conway's Game of Life


- World: a 2D grid. Each square represents a cell
- States: Living or dead
- Neighbourhood: The eight surrounding cells
- Initialize with a random number of living cells
- State transition rules:
 - ✱ A living cell with <2 living neighbours dies (loneliness)
 - ✱ A living cell with >3 living neighbours dies (overcrowded)
 - ✱ A dead cell with exactly 3 living neighbours comes alive
 - ✱ All other cells keep their current state





UPPSALA
UNIVERSITET

Life demo



Machine Learning course

- Spring semester (10 credits)
- Natural computation algorithms
 - ✱ all the above, and more
 - ✱ emphasis on neural networks
- Flexible choice of end project